Econometrics

Probability: 🎲, 🔮, 🪙

Florian Oswald
UniTo ESOMAS
2025-10-09

UNIVERSITÀ DI TORINO

Department of
Economics, Social Studies,
Applied Mathematics
and Statistics

Probability

This set of slides covers concepts from Stock and Watson chapter 2.

Probabilities

Probability is a function that assigns a value in $[0, 1]$ to a *set* (representing *events*).

Consider a fair 6-sided dice.

- Outcome space: $\Omega = \{1, 2, 3, 4, 5, 6\}$
- Event: a partition of Ω.

Probabilities

Probability is a function that assigns a value in $[0, 1]$ to a *set* (representing *events*).

Consider a fair 6-sided dice.

- Outcome space: $\Omega = \{1, 2, 3, 4, 5, 6\}$
- Event: a partition of Ω.

For example, the events A ("the outcome is even") and B ("the outcome is odd") are:

$$A = \{2, 4, 6\},$$
$$B = \{1, 3, 5\}.$$

Having a *fair* dice means that

$$\Pr(1) = \cdots = \Pr(6) = \frac{1}{6}$$

Task 1: Probabilities

Given a fair 6-sided dice , i.e. an outcome space $\Omega = \{1, 2, 3, 4, 5, 6\}$ and events

$$A = \{2, 4, 6\},$$
$$B = \{1, 3, 5\}.$$

how does one determine the ***likelihood*** (or the **probability**) of event A occuring?

Discrete Random Variables

A *Discrete Random Variable* is a function mapping outcomes to measurements.

For example, we might call

- X the number we obtain from throwing the dice once.
- X_1, X_2 the 2 numbers we obtain from throwing the dice two times.
- D an indicator (either `0`, `1`) for whether a randomly sampled person answers "yes" or "no" when asked whether they have children.

Gambler's Ruin

The dealer tosses a fair coin. If it comes up tails (`T`) you win, if it's heads (`H`)you loose.

Suppose we've seen the following sequence of tosses so far:

```
outcome:  H   H   H   H   H   H   H   H   H   H   ?
toss num: 1.  2.  3.  4.  5.  6.  7.  8.  9.  10. 11.
```

You have lost 10 times in a row now. Does this increase the probability that you will win the next round?

Gambler's Ruin

The dealer tosses a fair coin. If it comes up tails (`T`) you win, if it's heads (`H`)you loose.

Suppose we've seen the following sequence of tosses so far:

```
outcome:  H   H   H   H   H   H   H   H   H   H   ?
toss num: 1.  2.  3.  4.  5.  6.  7.  8.  9.  10. 11.
```

You have lost 10 times in a row now. Does this increase the probability that you will win the next round?

Which sequence is more likely to occur?

1. `TTTTTTTTTT` (10 `T` in a row)?
2. `HTHTTHTHHT` (random occurence of `T` and `H`)

Gambler's Ruin

The probability of hitting 11 `H` in a row is

$$\Pr(11 \text{ Heads in a Row}) = \frac{1}{2^{11}} = 0.00048 = 0.05\%$$

Gambler's Ruin

The probability of hitting 11 `H` in a row is

$$\Pr(\text{11 Heads in a Row}) = \frac{1}{2^{11}} = 0.00048 = 0.05\%$$

Ok, but I already had 10 `H`! What's the next toss likely going to be given that miserable history?

$$\Pr(\text{toss 11 is Heads} | \text{10 Heads before})$$

Let's calculate that probability on the board!

Task 2: 2 Dice

Given two fair 6-sided dice 🎲🎲, what is the probability of obtaining *at least* once the face "5"?

Probability Distributions

PDF/PMF and CDF of discrete RVs

- PDF/PMF: Probability Density/Mass Function. Table listing each outcome and the associated probability of observing it.
- CDF: Cumulative Distribution Function. Probability that a given RV takes on a value *less than or equal a certain value* (requires a notion of *ordering* - e.g. what about a dice with 6 different **colors**?)

For our 6-sided fair dice

```
##   outcome pdf cdf
## 1       1 1/6 1/6
## 2       2 1/6 2/6
## 3       3 1/6 3/6
## 4       4 1/6 4/6
## 5       5 1/6 5/6
## 6       6 1/6 6/6
```

PDF/PMF and CDF of discrete RVs

- PDF/PMF: Probability Density/Mass Function. Table listing each outcome and the associated probability of observing it.
- CDF: Cumulative Distribution Function. Probability that a given RV takes on a value *less than or equal a certain value* (requires a notion of *ordering* - e.g. what about a dice with 6 different **colors**?)

For our 6-sided fair dice

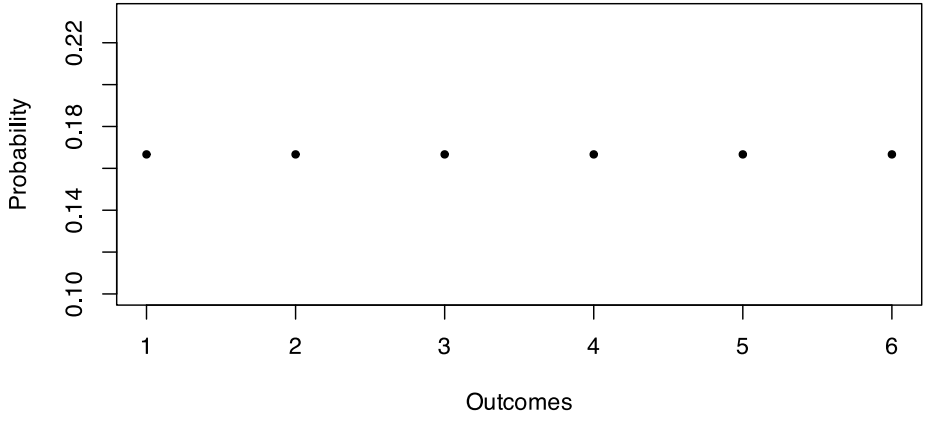

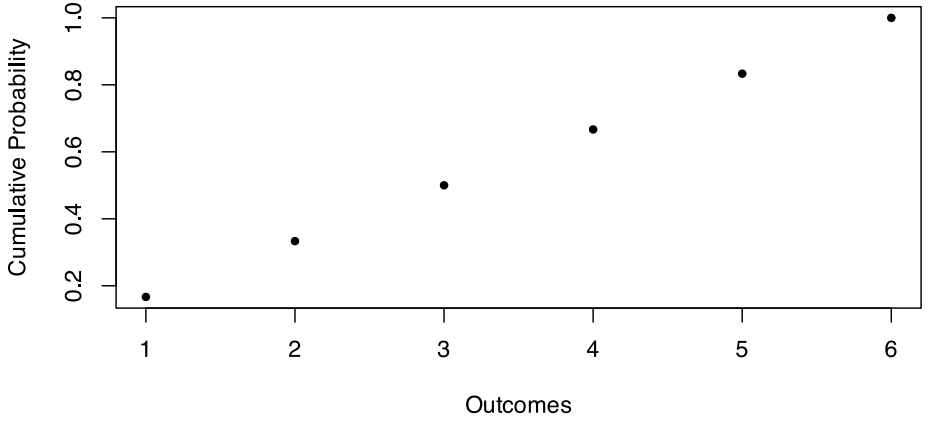

PDF/PMF and CDF of discrete RVs

- PDF/PMF: Probability Density/Mass Function. Table listing each outcome and the associated probability of observing it.
- CDF: Cumulative Distribution Function. Probability that a given RV takes on a value *less than or equal a certain value* (requires a notion of *ordering* - e.g. what about a dice with 6 different **colors**?)

For our fair coin

```
##   outcome pdf cdf
## 1    H(0) 1/2 1/2
## 2    T(1) 1/2 2/2
```

PDF/PMF and CDF of discrete RVs

- PDF/PMF: Probability Density/Mass Function. Table listing each outcome and the associated probability of observing it.
- CDF: Cumulative Distribution Function. Probability that a given RV takes on a value *less than or equal a certain value* (requires a notion of *ordering* - e.g. what about a dice with 6 different **colors**?)

For our fair coin

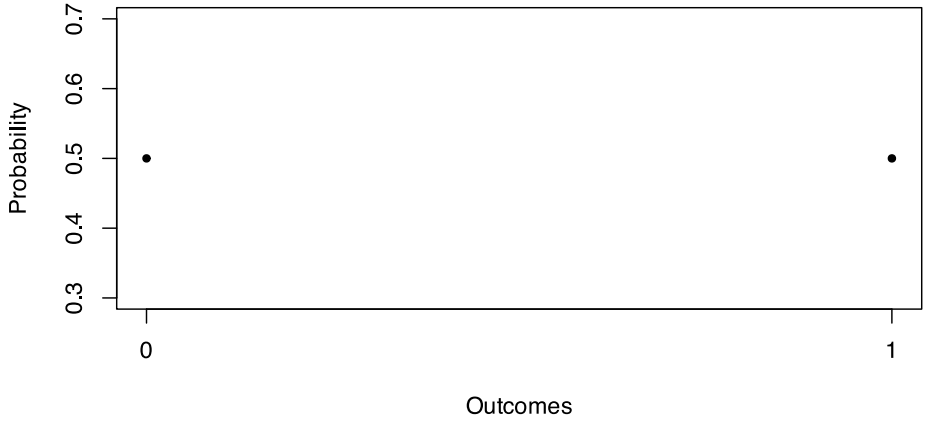

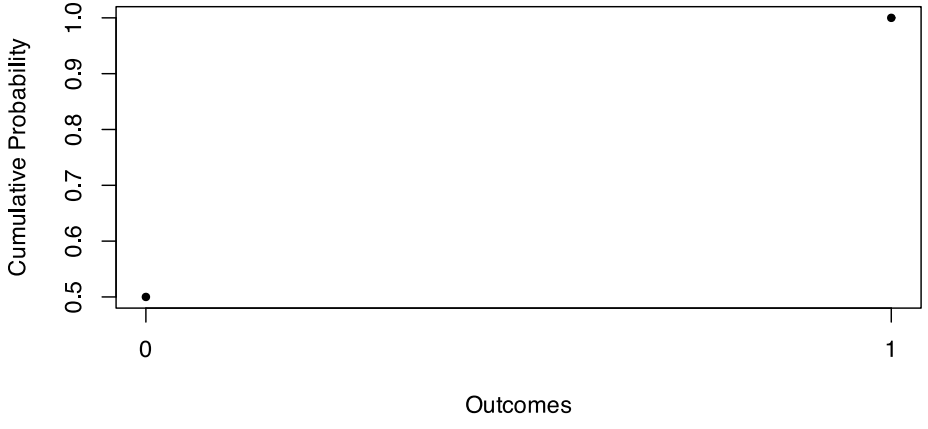

Bernoulli Distribution

- A special discrete RV with 2 outcomes: `0` and `1`, where event `1` ("success") occurs with probability p.

For example:

- Tomorrow it will rain with probability p (it will *not* rain with $1 - p$).

Bernoulli Distribution

- A special discrete RV with 2 outcomes: `0` and `1`, where event `1` ("success") occurs with probability p.

For example:

- Tomorrow it will rain with probability p (it will *not* rain with $1 - p$).

What kind of RV is flipping a *fair* coin? What about an *unfair* coin?

PDFs and CDFs of continuous Variables

- slightly more complicated because we need calculus and integration.
- The basic idea is the same!

Expected Value and Variance

Expected Value and Variance

We write

$$E(Y) = y_1 p_1 + \cdots + y_n p_n$$

For example for our dice example, where X is the result of throwing the dice:

$$\begin{aligned} E(X) &= 1 \times \Pr(1) + 2 \times \Pr(2) + \cdots + 6 \times \Pr(6) \\ &= \end{aligned}$$

Everybody calculate this result now!

Expected Value and Variance

We write

$$E(Y) = y_1 p_1 + \cdots + y_n p_n$$

For example for our dice example, where X is the result of throwing the dice:

$$\begin{aligned} E(X) &= 1 \times \Pr(1) + 2 \times \Pr(2) + \cdots + 6 \times \Pr(6) \\ &= 1 \times \frac{1}{6} + 2 \times \frac{1}{6} + \cdots + 6 \times \frac{1}{6} \\ &= 3.5 \end{aligned}$$

- So: We need to know ***all*** weights and values ($\Pr$ and X) in order to compute this quantity.

$E(X)$ is a Theoretical Concept. Mission Impossible?

- Imagine you get this message:

$E(X)$ is a Theoretical Concept. Mission Impossible?

- Imagine you get this message:

1. Your mission, should you choose to accept it, is to inspect this device : 🔮. Whenever you touch it, it displays a number in $\{1, \ldots, 10\}$. E.g. `4,8,4,9,1,5,10,9,1`...kind of random.
2. We must know the long-run average, i.e. E(🔮). Or something terrible will happen!
3. This message will destroy itself in 10,9,8,...💣

$E(X)$ is a Theoretical Concept. Mission Impossible?

- Imagine you get this message:

1. Your mission, should you choose to accept it, is to inspect this device : 🔮. Whenever you touch it, it displays a number in $\{1, \ldots, 10\}$. E.g. `4,8,4,9,1,5,10,9,1`...kind of random.
2. We must know the long-run average, i.e. E(🔮). Or something terrible will happen!
3. This message will destroy itself in 10,9,8,...💣

- You **don't know** the theoretical distribution of all the numbers in 🔮 (the $\Pr$'s). Time is running ⏰

$E(X)$ is a Theoretical Concept. Mission Impossible?

- Imagine you get this message:

1. Your mission, should you choose to accept it, is to inspect this device : 🔮. Whenever you touch it, it displays a number in $\{1, \ldots, 10\}$. E.g. `4,8,4,9,1,5,10,9,1`...kind of random.
2. We must know the long-run average, i.e. E(🔮). Or something terrible will happen!
3. This message will destroy itself in 10,9,8,...💣

- You **don't know** the theoretical distribution of all the numbers in 🔮 (the $\Pr$'s). Time is running ⏰

- What could you do? 🤔

TOM CRUISE
MISSION: IMPOSSIBLE
DEAD RECKONING

$E(X)$ is a Theoretical Concept. Mission Impossible?

- If we had a huge number of observations from 🔮 (a *sample*), we could come up with a *guess* of E(🔮), based on the data we got. Empirical, like.

```
set.seed(12345)  # to ensure reproducibility
Ps = runif(10)
Ps[c(2,3)] = 0
Ps = Ps / sum(Ps) # generate random weights
x = sample(1:10, size = 10000, replace = TRUE, prob = Ps)
xbar = mean(x)
xbar
```

```
## [1] 6.3418
```

$E(X)$ is a Theoretical Concept. Mission Impossible?

- If we had a huge number of observations from 🔮 (a *sample*), we could come up with a *guess* of E(🔮), based on the data we got. Empirical, like.

```
set.seed(12345)  # to ensure reproducibility
Ps = runif(10)
Ps[c(2,3)] = 0
Ps = Ps / sum(Ps) # generate random weights
x = sample(1:10, size = 10000, replace = TRUE, prob = Ps)
xbar = mean(x)
xbar
```

```
## [1] 6.3418
```

- `xbar` = $\bar{x} = \frac{1}{N}\sum_i^N x_i$ is called *Sample Mean*, *Arithmetic Average*, *Sample Average*
- $\bar{x}$ is an ***estimator*** for $E(X) = \mu_X$. ($E(X) = \mu_X$ by the way.)

Central Tendency - Mean and Median

`mean(x)`: the average of all values in `x`.

$$E(X) = \mu_X = \frac{1}{N}\sum_{i=1}^{N} x_i = \bar{x}$$

The second equality 👆 is correct only if...???

```
x <- c(1,2,2,2,2,100)
mean(x)
```

```
## [1] 18.16667
```

```
mean(x) == sum(x) / length(x)
```

```
## [1] FALSE
```

Central Tendency - Mean and Median

`mean(x)`: the average of all values in `x`.

$$E(X) = \mu_X = \frac{1}{N}\sum_{i=1}^{N} x_i = \bar{x}$$

The second equality 👆 is correct only if...???

```
x <- c(1,2,2,2,2,100)
mean(x)
```

```
## [1] 18.16667
```

```
mean(x) == sum(x) / length(x)
```

```
## [1] FALSE
```

`median`: the value x_j below and above which 50% of the values in `x` lie. m is the median if

$$\Pr(X \leq m) \geq 0.5 \text{ and } \Pr(X \geq m) \geq 0.5$$

The median is robust against *outliers*.

```
median(x)
```

```
## [1] 2
```

Quick Review

1. EV of a bernoulli
2. Continuous RV

Variance and Standard Deviation

- A measure of *spread* of a distribution.
- The definition of *variance* is

$$var(Y) = E[(Y - \mu_Y)^2]$$

if Y is discrete,

$$var(Y) = \sum_{i=1}^{N} (y_i - \mu_Y)^2 \times p_i.$$

Variance and Standard Deviation

- A measure of *spread* of a distribution.
- The definition of *variance* is

$$var(Y) = E[(Y - \mu_Y)^2]$$

if Y is discrete,

$$var(Y) = \sum_{i=1}^{N} (y_i - \mu_Y)^2 \times p_i.$$

- There is an issue with scaling: we *square* deviations.
- The standard deviation scales back to units of the data:

$$\sigma_Y = \sqrt{var(Y)}$$

Variance

Consider two `normal distributions` with equal mean at `0`:

Variance

Consider two `normal distributions` with equal mean at `0`:

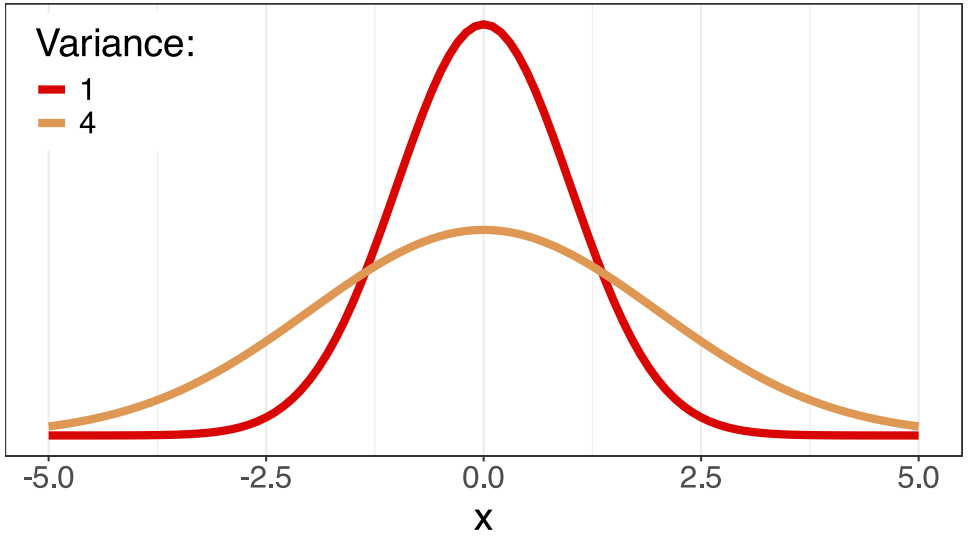

Compute with:

```
var(x)
```

Example of other Spread Measures

```r
# % catholic in 47 french-speaking
# swiss cantons in 1888
plot(swiss$Catholic,rep(1,nrow(swiss)),pch = 3,
     cex = 2,xlab = "% Catholic",yaxt = "n",ylab = "")
```

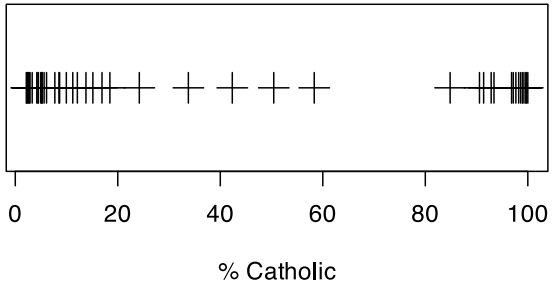

How do the values in column `Catholic` *vary*?

Measure	R	Result
Variance	`var(swiss$Catholic)`	1739.29
Standard Deviation	`sd(swiss$Catholic)`	41.7
IQR	`IQR(swiss$Catholic)`	87.93
Minimum	`min(swiss$Catholic)`	2.15
Maximum	`max(swiss$Catholic)`	100
Range	`range(swiss$Catholic)`	2.15, 100

Task 3: Computing Variance by hand

1. Compute the variance of our 6-sided dice!
2. compute the variance of $X \sim \text{Bernoulli}(p)$! Bonus question: what value of p maximizes this variance?

Two Random Variables

(X, Y), (,), (,), (,)

Two Random Variables: Example

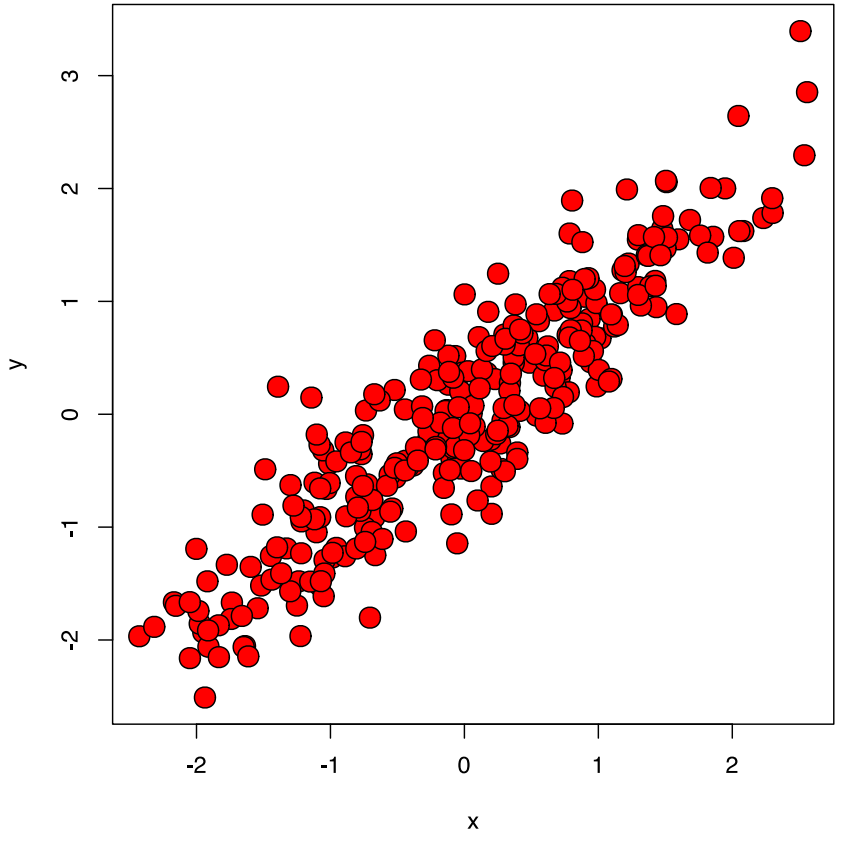

- Here, X and Y are ***joint normally*** distributed.
- We would write

$$(X, Y) \sim \mathcal{N}\Big([\, \mu_X \; \mu_Y \,] \,,\, \Sigma \Big)$$

`

where Σ is a *matrix*

$$\begin{bmatrix} \sigma_X^2 & \rho\, \sigma_X \sigma_Y \\ \rho\, \sigma_X \sigma_Y & \sigma_Y^2 \end{bmatrix}$$

Tabulating Data

`table(x)` is a useful function that counts the occurence of each unique value in `x`:

```
# install.packages("dslabs") in order to use this command
data("gapminder",package = "dslabs")
table(gapminder$continent)
```

```
##
##   Africa Americas     Asia   Europe  Oceania
##     2907     2052     2679     2223      684
```

Tabulating Data

`table(x)` is a useful function that counts the occurence of each unique value in `x`:

```
# install.packages("dslabs") in order to use this command
data("gapminder",package = "dslabs")
table(gapminder$continent)
```

```
##
##   Africa Americas     Asia   Europe  Oceania
##     2907     2052     2679     2223      684
```

The same can be achieved using the `count` function (from `dplyr`)

```
gapminder %>% count(continent)
```

```
##   continent    n
## 1    Africa 2907
## 2  Americas 2052
## 3      Asia 2679
## 4    Europe 2223
## 5   Oceania  684
```

Tabulating Data

Given two variables, `table` produces a contingency table:

```
gapminder_new <- gapminder %>%
  filter(year == 2015) %>%
  mutate(fertility_above_2 = (fertility > 2.1)) # dummy variable for fertility rate above replacement rate
```

```
table(gapminder_new$fertility_above_2)
```

```
##
## FALSE  TRUE
##    80   104
```

```
table(gapminder_new$continent,gapminder_new$fertility_above_2)
```

```
##
##            FALSE TRUE
##   Africa       2   49
##   Americas    15   20
##   Asia        20   27
##   Europe      39    0
##   Oceania      4    8
```

Cross-Tabulating Data: A joint distribution!

The probability that $X = x$ and $Y = y$ is

$$\Pr(X = x, Y = y)$$

```
# table with absolute numbers
abstab = table(gapminder_new$continent,
               gapminder_new$fertility_above_2)
(jd <- round(prop.table(abstab),2))  # proportions in
```

```
##
##            FALSE TRUE
##   Africa    0.01 0.27
##   Americas  0.08 0.11
##   Asia      0.11 0.15
##   Europe    0.21 0.00
##   Oceania   0.02 0.04
```

- check: proper probability distribution?

```
sum(jd)  # sums up all elements
```

```
## [1] 1
```

Cross-Tabulating Data: A joint distribution!

The probability that $X = x$ and $Y = y$ is

$$\Pr(X = x, Y = y)$$

```
# table with absolute numbers
abstab = table(gapminder_new$continent,
               gapminder_new$fertility_above_2)
(jd <- round(prop.table(abstab),2))  # proportions in
```

```
##
##            FALSE TRUE
##   Africa    0.01 0.27
##   Americas  0.08 0.11
##   Asia      0.11 0.15
##   Europe    0.21 0.00
##   Oceania   0.02 0.04
```

- check: proper probability distribution?

```
sum(jd)  # sums up all elements
```

```
## [1] 1
```

- for instance

$$\Pr(\text{Asia}, \text{TRUE}) = 0.14673$$

- You can see 👀 that this is just the **share** of cases in each bin!

Tabulating Data: marginal distributions!

```r
# proportions by row
(m1 = prop.table(abstab, margin = 1))
```

```
##
##                 FALSE       TRUE
##   Africa   0.03921569 0.96078431
##   Americas 0.42857143 0.57142857
##   Asia     0.42553191 0.57446809
##   Europe   1.00000000 0.00000000
##   Oceania  0.33333333 0.66666667
```

check?

```r
rowSums(m1)
```

```
##   Africa Americas     Asia   Europe  Oceania
##        1        1        1        1        1
```

Tabulating Data: marginal distributions!

```r
# proportions by row
(m1 = prop.table(abstab, margin = 1))
```

```
##
##                 FALSE       TRUE
##   Africa   0.03921569 0.96078431
##   Americas 0.42857143 0.57142857
##   Asia     0.42553191 0.57446809
##   Europe   1.00000000 0.00000000
##   Oceania  0.33333333 0.66666667
```

check?

```r
rowSums(m1)
```

```
##   Africa Americas     Asia   Europe  Oceania
##        1        1        1        1        1
```

```r
# proportions by column
(m2 = prop.table(abstab, margin = 2))
```

```
##
##                 FALSE       TRUE
##   Africa   0.02500000 0.47115385
##   Americas 0.18750000 0.19230769
##   Asia     0.25000000 0.25961538
##   Europe   0.48750000 0.00000000
##   Oceania  0.05000000 0.07692308
```

check?

```r
colSums(m2)
```

```
## FALSE  TRUE
##     1     1
```

Joint Distributions

The probability that $X = x$ and $Y = y$ is

$$\Pr(X = x, Y = y)$$

```
##                      Rain (X=0) No Rain (X=1)
## Long Commute (Y=0)        0.15          0.07
## Short Commute (Y=1)       0.15          0.63
```

for instance,

- $\Pr(Y = 1, X = 0) = 0.15$

Marginal Distributions

$$\Pr(Y = y) = \sum_{i=1}^{k} \Pr(X = x_i, Y = y)$$

```
##                      Rain (X=0) No Rain (X=1)
## Long Commute (Y=0)        0.15          0.07
## Short Commute (Y=1)       0.15          0.63
```

- Probability of long commute, *irrespective of rain*, is 0.15 + 0.07 = 0.22
- Probability of rain, *irrespective of length of commute*, is 0.15 + 0.15 = 0.3

Task 4 : Computing Marginal Distributions

05:00

```
r
```

```
##                      Rain (X=0) No Rain (X=1)
## Long Commute (Y=0)        0.15          0.07
## Short Commute (Y=1)       0.15          0.63
```

1. Compute the marginal distribution of X
2. Compute the marginal distribution of Y

Conditional distributions

- Conditional distribution is $\Pr(X = x | Y = y) = \frac{\Pr(X=x, Y=y)}{\Pr(Y=y)}$
- We **fix** the value of one variable, and look at the resulting distribution of the other variable.
- Notice, the result must be a *valid* probability distribution, i.e., it must sum to 1.
- We compute the probability as before, but restrict attention to where the condition we impose is true.

Task 5: compute conditional distribution

03:00

Consider

```
##                      Rain (X=0) No Rain (X=1)
## Long Commute (Y=0)        0.15          0.07
## Short Commute (Y=1)       0.15          0.63
```

- Suppose we know that it does not rain $x = 1$. What is the distribution of Y, *given* this knowledge, i.e. what is the distribution $\Pr(Y|X = 1)$?

Task 6: Computing More Conditional Distributions

05:00

```
jd
```

```
##
##             FALSE TRUE
##   Africa     0.01 0.27
##   Americas   0.08 0.11
##   Asia       0.11 0.15
##   Europe     0.21 0.00
##   Oceania    0.02 0.04
```

1. Compute the Distribution of countries conditional on low fertility?
2. Compute the Distribution of countries conditional on being in Asia?
3. Compute the marginal distribution of high/low fertility

Conditional Expectation

We define

$$E(Y|X=x) = \sum_{i=1}^{k} y_i \Pr(Y = y_i | X = x)$$

- Notice the y_i there!
- Basically this is expected value of Y, but **under the assumption that** X takes on the value x.
- We *fix* the joint distribution of (Y, X) at a certain value $X = x$.

Task 7: Conditional Expectation

TABLE 2.3 Joint and Conditional Distributions of Number of Wireless Connection Failures (M) and Network Age (A)

A. Joint Distribution

	$M = 0$	$M = 1$	$M = 2$	$M = 3$	$M = 4$	Total
Old network ($A = 0$)	0.35	0.065	0.05	0.025	0.01	0.50
New network ($A = 1$)	0.45	0.035	0.01	0.005	0.00	0.50
Total	0.80	0.10	0.06	0.03	0.01	1.00

B. Conditional Distributions of M given A

	$M = 0$	$M = 1$	$M = 2$	$M = 3$	$M = 4$	Total
$\Pr(M \mid A = 0)$	0.70	0.13	0.10	0.05	0.02	1.00
$\Pr(M \mid A = 1)$	0.90	0.07	0.02	0.01	0.00	1.00

1. Compute $E(M)$
2. Compute $E(M|A = 0)$

Law of Iterated Expectations

- We can decompose the expected value of Y into subgroups, see what fraction of total probability each group makes up, and compute their weighted average.

Law of Iterated Expectations

- We can decompose the expected value of Y into subgroups, see what fraction of total probability each group makes up, and compute their weighted average.

```
##    name sex height
## 1 Peter   M   1.90
## 2 James   M   1.75
## 3  Mary   F   1.68
## 4  John   M   1.86
## 5   Amy   F   1.59
```

- The average height here is $(1.9 + 1.75 + 1.68 + 1.86 + 1.59)/5 = 1.756$:

Law of Iterated Expectations

> The average height here is 1.756

- What is the average for each **group** by sex?

```
## # A tibble: 2 × 3
##   sex       n mean_height
##   <chr> <int>       <dbl>
## 1 F         2        1.64
## 2 M         3        1.84
```

Law of Iterated Expectations

> The average height here is 1.756

- What is the average for each **group** by sex?

```
## # A tibble: 2 × 3
##   sex       n mean_height
##   <chr> <int>       <dbl>
## 1 F         2        1.64
## 2 M         3        1.84
```

- Can we recover the unconditional mean from this information? Yes we can!
- We need again *weights* and *values*, like for the standard $E(Y)$

```r
weights = c(2/5, 3/5)
values = c(1.635, 1.836667)
sum(weights * values) # 2/5 * 1.635 + 3/5 * 1.836667
```

```
## [1] 1.756
```

Law of Iterated Expectations

- What did we just do?
- We established that we can get the unconditional mean of Y from the *conditional* one, if we know the *weights* we have to give to each subgroup. More precisely:

1. We know $E(Y|X)$:
    1. $E(Y|\text{male}) = 1.836667$
    2. $E(Y|\text{female}) = 1.635$
2. We know the proportion of both groups:
    1. $\Pr(\text{male}) = 3/5$
    2. $\Pr(\text{female}) = 2/5$

Therefore, we can recompose $E(Y)$:

$$\begin{aligned} E(Y) &= \sum_{i=1}^{k} E(Y|X = x_i) \Pr(X = x_i) \\ &= E_X \left[E(Y|X) \right] \end{aligned}$$

Law of Iterated Expectations

- One particular version of this comes up often in econometrics:

$$E[Y|X] = 0 \Rightarrow E[Y] = 0$$

- Why? Well, now you know that $E[Y|X] = 0$ for all possible X is zero.
- Compute $E(Y) = \sum_{i=1}^{k} 0 \times \Pr(X = x_i)$!

Conditional Expectation and Prediction Errors

- We said that the the squared error makes sense (talking about variance)
- Suppose we want to design a **predictor** $g(X)$ which can **predict** Y reasonably well. One criterion for that might be the *average squared error of prediction*.

Conditional Expectation and Prediction Errors

- We said that the the squared error makes sense (talking about variance)
- Suppose we want to design a **predictor** $g(X)$ which can **predict** Y reasonably well. One criterion for that might be the *average squared error of prediction*.
- In other words, what is the solution to this problem?

$$\min_m E\left[(Y-m)^2\right]$$

Independence

- We call 2 RVs *independent* if knowing the value of one of them is not informative about the value of the other.
- The precise definition of this statement is

$$\Pr(X = x | Y = y) = \Pr(X = x)$$

Independence

- We call 2 RVs *independent* if knowing the value of one of them is not informative about the value of the other.
- The precise definition of this statement is

$$\Pr(X = x | Y = y) = \Pr(X = x)$$

- or:

$$\Pr(X = x, Y = y) = \Pr(X = x) \Pr(Y = y)$$

Covariance

- Covariance tells us how two RVs co-vary.

$$
\begin{aligned}
cov(X, Y) = \sigma_{XY} &= E\left[(X - \mu_X)(Y - \mu_Y)\right] \\
&= \sum_{i=1}^{k} \sum_{j=1}^{m} (x_i - \mu_X)(y_j - \mu_Y) \Pr(X = x_i, Y = y_j)
\end{aligned}
$$

- Notice how this object tells us how large deviations of X from its mean *jointly* occur with such deviations of Y.

- This confounds the scale of X and Y, however, so not easy to interpret

Correlation

- Correlation fixes this issue.

$$corr(X,Y) = \frac{cov(X,Y)}{\sqrt{var(X)var(Y)}} = \frac{\sigma_{XY}}{\sigma_X \sigma_Y}$$

- Correlation is always between -1 and 1.
- $corr(X,Y) = 0$ means they are *uncorrelated*.

Correlation and Conditional Mean

- If the conditional mean of Y does **not** depend on X, then they must be uncorrelated.
- That is, if

$$E(Y|X) = \mu_Y \Rightarrow cov(X, Y) = 0 \text{ and } corr(X, Y) = 0$$

Task 8 : Show Correlation and Conditional Mean

Show this result!

$$E(Y|X) = \mu_Y \Rightarrow cov(X, Y) = 0 \text{ and } corr(X, Y) = 0$$

Random Sampling (SW KC 2.5)

- Suppose we **randomly** select n objects from a large population.
- Y_i is the value of i-th object chosen.
- Each object i has the same chance of getting drawn, which means that the *distribution of the Ys is the same*, and object i being drawn does not depend on object j being drawn.
- **Same** means **identical**. So, the Ys are ***identically and independently*** distributed.

Sampling Distribution of Sample Average

- We already defined the **sample average**:

$$\bar{y} = \frac{1}{n}(y_1 + y_2 + \cdots + y_n) = \frac{1}{n}\sum_{i=1}^{n} y_i$$

Sampling Distribution of Sample Average

- We already defined the **sample average**:

$$\bar{y} = \frac{1}{n}(y_1 + y_2 + \cdots + y_n) = \frac{1}{n}\sum_{i=1}^{n} y_i$$

- A core concept is to realize that the value of $\bar{y}$ is itself ***random***. Just imagine drawing a second sample of y's

Sampling Distribution of Sample Average

- We already defined the **sample average**:

$$\bar{y} = \frac{1}{n}(y_1 + y_2 + \cdots + y_n) = \frac{1}{n}\sum_{i=1}^{n} y_i$$

- A core concept is to realize that the value of $\bar{y}$ is itself ***random***. Just imagine drawing a second sample of y's
- Because $\bar{y}$ varies with each sample, the object $\bar{y}$ has a **sampling distribution**.
- Every well-defined distribution has a **mean** and a **variance** - same for this sampling distribution. Let's figure it out!

Mean of the Sample Average

- The mean of the sample average (*the mean of the mean!*) is just

$$E(\bar{y}) = E\left[\frac{1}{n}\sum_{i=1}^{n} y_i\right] = \frac{1}{n}\sum_{i=1}^{n} E[y_i] = \frac{1}{n} n\mu_Y$$

Mean of the Sample Average

- The mean of the sample average (*the mean of the mean!*) is just

$$E(\bar{y}) = E\left[\frac{1}{n}\sum_{i=1}^{n} y_i\right] = \frac{1}{n}\sum_{i=1}^{n} E[y_i] = \frac{1}{n} n\mu_Y$$

- So that's good news. The expected value is equal to the population mean.
- This statement is important. We say that $\bar{y}$ is an ***unbiased estimator*** for $E(y)$.

Task 9: Variance of the Sample Average

- Similarly,

$$var(\bar{y}) = var\left(\frac{1}{n}\sum_{i=1}^{n} y_i\right) = \frac{\sigma_Y^2}{n}$$

- Show this result!

Properties of the Sample Mean

The following facts **always** hold for the sample mean:

1. $E(\bar{y}) = \mu_Y$
2. $var(\bar{y}) = \frac{\sigma_Y^2}{n}$
3. $std.dev(\bar{y}) = \frac{\sigma_Y}{\sqrt{n}}$

Properties of the Sample Mean

The following facts **always** hold for the sample mean:

1. $E(\bar{y}) = \mu_Y$
2. $var(\bar{y}) = \frac{\sigma_Y^2}{n}$
3. $std.dev(\bar{y}) = \frac{\sigma_Y}{\sqrt{n}}$

Those are true ***whenever y is i.i.d***. In particular, we do **not** need the y to be normally distributed!

Sampling from a Normal Distribution

1. The sum of normally distributed RVs is itself normally distributed.

Sampling from a Normal Distribution

1. The sum of normally distributed RVs is itself normally distributed.
2. Suppose we draw Y_i i.i.d from a normal distribution.

Sampling from a Normal Distribution

1. The sum of normally distributed RVs is itself normally distributed.
2. Suppose we draw Y_i i.i.d from a normal distribution.
3. We already know the mean and variance of the sample mean $\bar{y}$ in this case (previous slide.)

Sampling from a Normal Distribution

1. The sum of normally distributed RVs is itself normally distributed.
2. Suppose we draw Y_i i.i.d from a normal distribution.
3. We already know the mean and variance of the sample mean $\bar{y}$ in this case (previous slide.)
4. In this **special case** we actually ***know*** the full distribution of $\bar{y}$, which has to be *normal* itself (see point 1). That is

$$\bar{y} \sim \mathcal{N}\left(\mu_y, \frac{\sigma_Y^2}{n}\right)$$

Sampling from a Normal Distribution

1. The sum of normally distributed RVs is itself normally distributed.
2. Suppose we draw Y_i i.i.d from a normal distribution.
3. We already know the mean and variance of the sample mean $\bar{y}$ in this case (previous slide.)
4. In this **special case** we actually ***know*** the full distribution of $\bar{y}$, which has to be *normal* itself (see point 1). That is

$$\bar{y} \sim \mathcal{N}\left(\mu_y, \frac{\sigma_Y^2}{n}\right)$$

What if we **don't know** the distribution of our Y_i?

NOT Sampling from a Normal Distribution

- This is clearly the empirically relevant case. We almost **never** know the precise distribution of our data.
- Let us repeat that:

> The following facts **always** hold for the sample mean:
>
> 1. $E(\bar{y}) = \mu_Y$
> 2. $var(\bar{y}) = \frac{\sigma_Y^2}{n}$
> 3. $std.dev(\bar{y}) = \frac{\sigma_Y}{\sqrt{n}}$
>
> Those are true ***whenever y is i.i.d***. In particular, we do **not** need the y to be normally distributed!

Asymptotic Distributions

- If we don't know the *exact* distribution of the data, we can still derive an **approximate** distribution.
- The approximation requires that we have a lot of data, i.e. n is large.
- We are using 2 key tools for this: The Law of Large Numbers and the Central Limit Theorem.

The Law of Large Numbers

- We say that a sequence of RVs z_n **converges in probability** to the constant a if the probability that z_n comes to lie in the interval $[a - c, a + c]$ goes to 1 as n goes to infinity.
- We would write that as

$$z_n \xrightarrow{p} a$$

The Law of Large Numbers

- We say that a sequence of RVs z_n **converges in probability** to the constant a if the probability that z_n comes to lie in the interval $[a - c, a + c]$ goes to 1 as n goes to infinity.
- We would write that as

$$z_n \xrightarrow{p} a$$

The Law of Large Numbers says that *if* Y_i are i.i.d. with $E(Y) = \mu_Y$, then - if large outliers are very unlikely (if $var((Y)\sigma_Y^2 < \infty)$) - the sample mean converges in probability to the population mean, i.e.

$$\bar{y} \xrightarrow{p} \mu_Y$$

- This property ($\bar{y} \xrightarrow{p} \mu_Y$) is called ***consistency***: $\bar{y}$ is **consistent** for μ_Y.

Demo

```r
set.seed(45)
n = 2
p = 0.78
s = sample(c(0,1), size = n, prob = c(1-p, p), replace = TRUE)
print(s)
```

```
## [1] 1 1
```

```r
mean(s) # ybar
```

```
## [1] 1
```

```r
n = 10
s = sample(c(0,1), size = n, prob = c(1-p, p), replace = TRUE)
print(s)
```

```
##  [1] 1 1 1 1 1 1 1 1 1 0
```

```r
mean(s) # ybar
```

```
## [1] 0.9
```

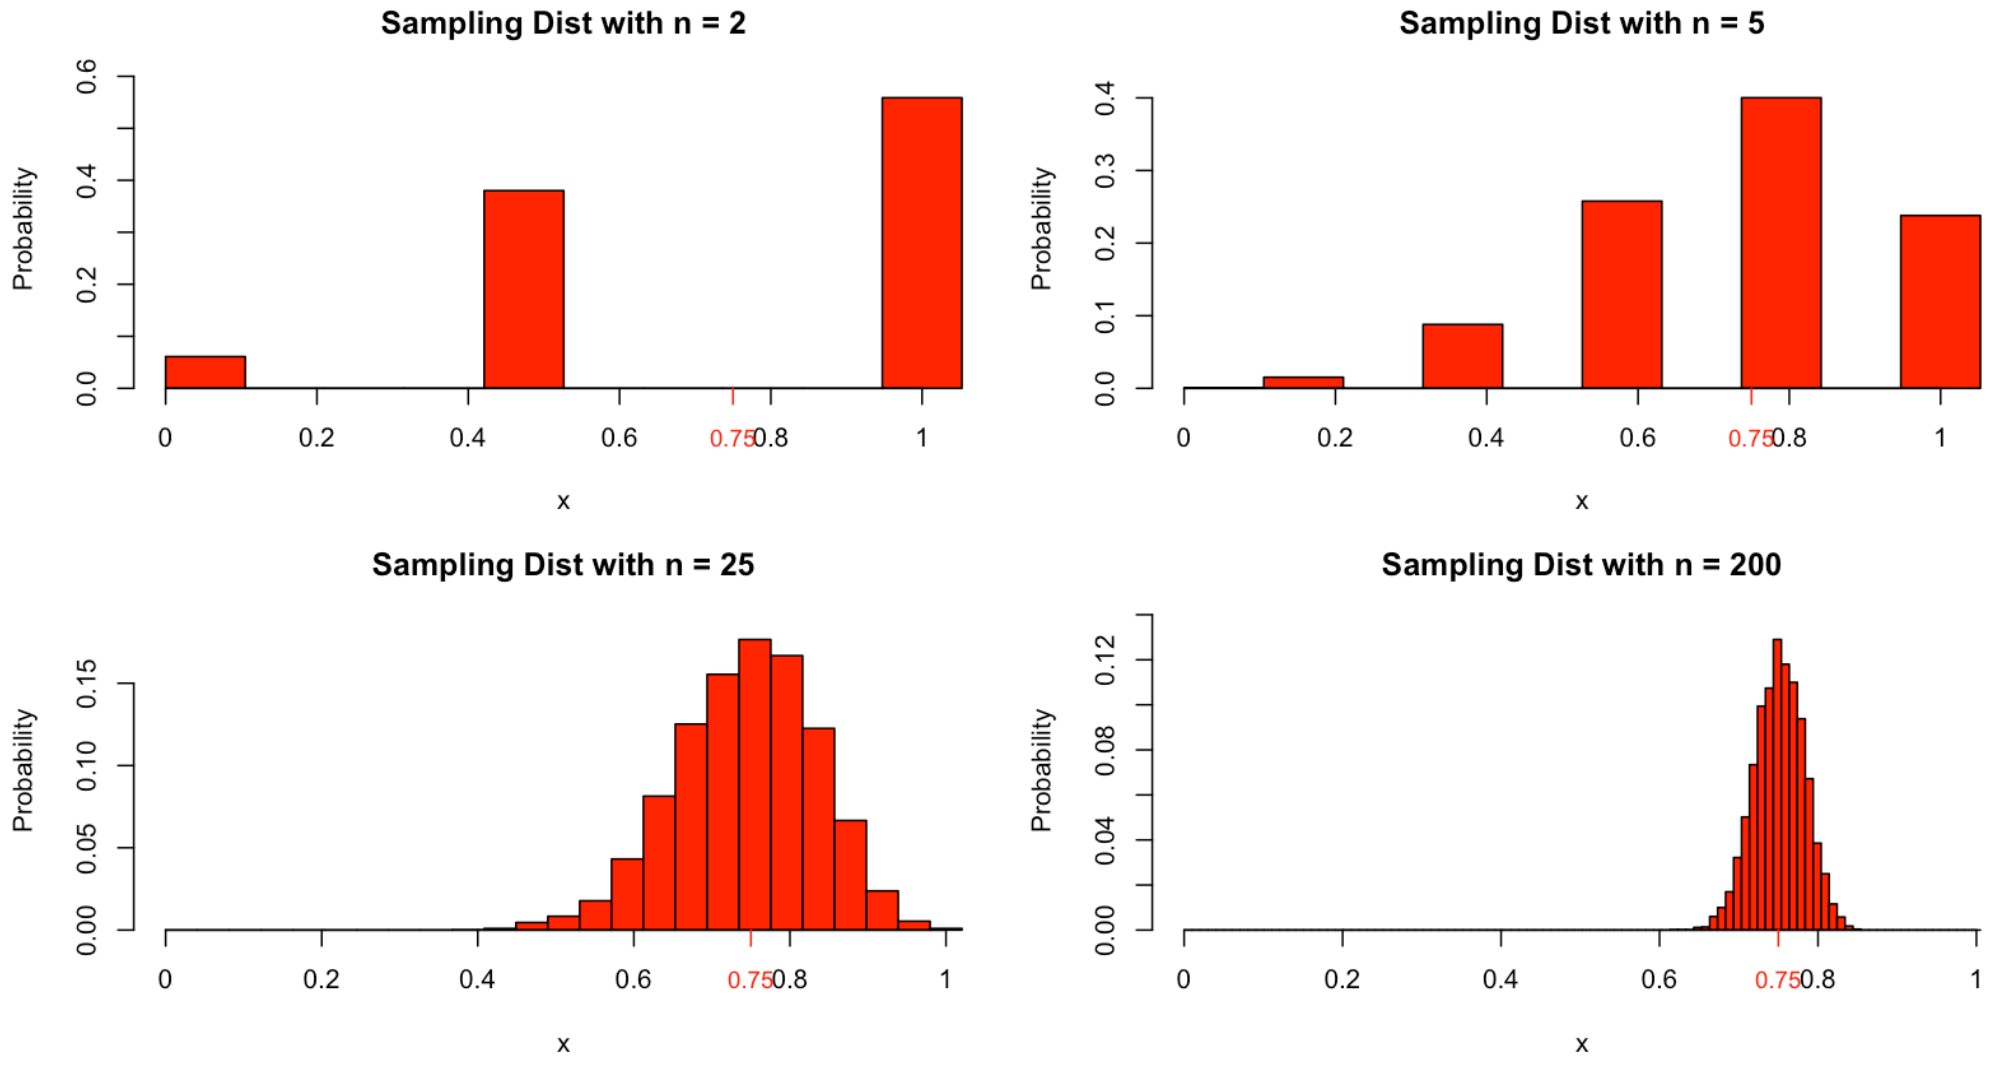
Sampling Dist with n = 2
Probability
0.0
0.2
0.4
0.6
0
0.2
0.4
0.6
0.75
0.8
1
x
Sampling Dist with n = 5
Probability
0.0
0.1
0.2
0.3
0.4
0
0.2
0.4
0.6
0.75
0.8
1
x
Sampling Dist with n = 25
Probability
0.00
0.05
0.10
0.15
0
0.2
0.4
0.6
0.75
0.8
1
x
Sampling Dist with n = 200
Probability
0.00
0.04
0.08
0.12
0
0.2
0.4
0.6
0.75
0.8
1
x

The Central Limit Theorem

- This important theorem says that under relavitely mild conditions, the distibution of $\bar{y}$ is **normal** when n is large.

- We already knew from above that the mean and variance of $\bar{y}$ are μ_Y and $\frac{\sigma_Y^2}{n}$ - *whatever* the underlying distribution of y was.

--

- But that $\bar{y}$ would converge to a normal if n large, that is new!

The Central Limit Theorem

- This important theorem says that under relavitely mild conditions, the distibution of $\bar{y}$ is **normal** when n is large.

- We already knew from above that the mean and variance of $\bar{y}$ are μ_Y and $\frac{\sigma_Y^2}{n}$ - *whatever* the underlying distribution of y was.

--

- But that $\bar{y}$ would converge to a normal if n large, that is new!

- This holds for ***any*** distribution that Y might have!

Demo

```r
library(shinyCLT)
CLT()
```

```
library(CLT)
CLT()
```

10:00

class: inverse

Task 4: Summarising data

1. Compute the mean of GDP in 2011 and assign to object `mean`. You should exclude missing values. (*Hint: read the help for* `mean` *to remove* `NA`*s*).

2. Compute the median of GDP in 2011 and assign to object `median`. Again, you should exclude missing values. Is it greater or smaller than the average?

3. Create a density plot of GDP in 2011 using `geom_density`. A density plot is a way of representing the distribution of a numeric variable. Add the following code to your plot to show the median and mean as vertical lines. What do you observe?
   ```
   geom_vline(xintercept = as.numeric(mean), colour = "red") +
   geom_vline(xintercept = as.numeric(median), colour = "orange")
   ```

4. Compute the correlation between fertility and infant mortality in 2015. To drop `NA`s in either variable set the argument `use` to "pairwise.complete.obs" in your `cor()` function. Is this correlation consistent with the graph you produced in Task 3?

Department of
Economics, Social Studies,
Applied Mathematics
and Statistics

UNIVERSITÀ
DI TORINO

That's all for Probability!

	Slides
	My Homepage
	Book